

# Project Report: Anime Face Generation using GANs

## Objective

The goal of this project was to design and train a **Generative Adversarial Network (GAN)** capable of generating high-quality **anime face images**. Using the publicly available Anime Face Dataset from Kaggle, the model was trained to learn the underlying patterns of anime-style facial features and synthesize novel, realistic-looking faces.

---

## Methodology

### 1. Data Preparation

- **Source:** Anime Face Dataset (~63,000 images)
  - **Preprocessing Steps:**
    - Loaded and decoded `.jpg` files from dataset folder.
    - Resized all images to **64×64 pixels** for efficient training.
    - Normalized pixel values to range **[-1, 1]** (as expected by `tanh` output layer).
  - **Batching:** Used TensorFlow's `tf.data.Dataset` API with `shuffle`, `batch (size = 256)`, and `prefetch` for optimized data pipeline.
- 

### 2. Model Architecture

#### Generator:

- **Input:** 100-dimensional random noise vector.
- **Layers:**
  - Dense → Reshape to 8×8×256
  - 3× Conv2DTranspose layers to upsample to 64×64
- **Activations:** ReLU (hidden layers), Tanh (output)

#### Discriminator:

- **Input:** 64×64×3 RGB image
- **Layers:**
  - 3× Conv2D layers with LeakyReLU and Dropout
  - Flatten → Dense
- **Output:** Single logit score (real/fake)

#### Losses:

- **Generator:** Binary Crossentropy (tries to fool discriminator)
  - **Discriminator:** Binary Crossentropy (classify real/fake)
- 

### 3. Training Strategy

- **Epochs:** 50
  - **Optimizers:** Adam with learning rate = 0.0001 and  $\beta_1 = 0.5$
  - **Batch Size:** 256
  - **Regular Checkpoints:** Model saved every 5 epochs
  - **Image Generation:** Snapshot of 16 generated images saved every epoch
- 

#### Qualitative Analysis:

- **Image Quality:**
  - Images are visually sharp with coherent anime facial structures.
- **Diversity:**
  - Generated faces show a variety of hairstyles, colors, and expressions.

### Challenges Faced

1. **Training Instability:**
  - Early epochs saw frequent mode collapse and poor discriminator loss.
  - Solution: Used `label smoothing` and tuned learning rates.
2. **Hardware Constraints:**
  - Training on Google Colab with limited VRAM required tuning batch sizes.
3. **Evaluation Difficulties:**
  - No labeled classes made using classifier-based metrics tricky.
  - Solution: Used pre-trained InceptionNet and resized images for FID/IS.
4. **GAN Saturation:**
  - At times, the discriminator became too strong.
  - Introduced Dropout and moderate label noise to improve balance.